

**COURSE WORK TITLE:**

**DATASET ANALYSIS AND VISUALISATION USING BIG DATA PROGRAM**

**COURSE WORK TOPIC:**

**PREDICTING AIRLINE CUSTOMER'S RECOMMENDATION THROUGH DATA  
INSIGHT**

**COLLEGE OF COMPUTING, ENVIRONMENT AND SCIENCE  
COVENTRY UNIVERSITY**

**NAME:**

**OBIANUJU VIVIAN OKAFOR**

**STUDENT ID:**

**14281642**

**COURSE:**

**BIG DATA ANALYTICS AND DATA VISUALISATION**

**APRIL 2024**

## **ABSTRACT**

Many people travel by air daily, and there are several airline businesses that are prepared to accommodate these consumers' needs. Customers do, however, select airlines based on the services they receive, and if they are happy with the airline and the services rendered, they also recommend the airline to friends and family.

These reviews have the potential to help these airlines provide better service.

The project's objective is to use big data analytics to anticipate, based on various airline service ratings, whether a customer's recommendation is good or negative. The Airline Review Dataset from Kaggle will be used for this investigation.

The dataset will undergo preprocessing and data cleaning, which will prepare the data for further modelling. Three different machine learning models—logistic regression, decision trees, and random forests—will be used in this study to generate predictions.

**Keywords:** Random forest, decision tree, logistic regression, and airline review dataset.

## TABLE OF CONTENT

Abstract.....	2
Introduction.....	4
Background/Related Work.....	5
Dataset Description.....	6
Methodology.....	7
Logistic Regression.....	8
Decision Tree.....	8
Random Forest.....	8
Software Installation.....	9
Experimental Section.....	9
Tableau Visualization.....	11
Result Discussion.....	12
Conclusion and Future Works.....	13
Social and Ethical Impact .....	13
References .....	14
Appendix I.....	15
Appendix II.....	30

## INTRODUCTION

With so many brands to choose from, consumers have high expectations. From the first contact to the after-sale period and beyond, every stage of the customer experience must involve ensuring their satisfaction.

Customer service is the help and direction a business offers to customers before, during, and after they make a purchase of a good or service.

Research shows that building client loyalty and promoting business growth depends on providing excellent customer service (Savage, 2024).

Customers' service preferences during interactions vary based on factors such as the nature of assistance required, and their level of personal familiarity with the product or service being evaluated for purchase. Many technologically savvy customers explore many information sources on the internet and frequently know far more about the various components of the interesting product or service than friendly representatives in the store (Lee & Lee, 2019).

Excellent customer service can improve a business's relationships with its customers and also help improve a business as customers tend to recommend a business based on their experience with the business and the services the business rendered.

To predict if a customer will likely recommend an airline based on their services the airline review dataset will be analyzed using big data analytics. The project uses machine learning techniques including Logistic Regression, Random Forest, and Decision Tree to model and predict customer's recommendation is positive or negative. Accuracy is used to evaluate how well these models perform at making reliable and accurate predictions. This project's software tools include Tableau, Jupyter Environment, and Pyspark. PySpark is our major analytics engine. Tableau helps visualise and explain patterns, resulting in analytically sound and meaningful solutions.

In summary, we will use PySpark and Tableau to conduct the analysis for this project.

## BACKGROUND/RELATED WORK

Patel et al. (2023) wrote that customer feedback is important as it is an opportunity for businesses to improve the services given to its customers. The study used sentimental analysis to analyze the airline reviews dataset. The study tested the performance of sentimental analysis using machine learning algorithms like Naïve Bayes, Support Vector Machine and Decision Tree.

Analysing customer experience and emotion creates a sustainable technology for everyone. Sustainable technology enables us to adopt factors such as stronger business output, a smoother user experience, attracting more customers etc.

To thrive in this highly competitive environment, airlines are constantly working to improve service quality, which has become a key survival strategy. Just as product quality transformed manufacturing competitiveness, the success of the aviation sector now depends greatly on creative service quality, which has become the distinguishing factor between an airline's success and failure. Continuous improvement in service quality is recognised as a critical strategy for creating a competitive advantage and ensuring customer satisfaction. However, many airlines continue to struggle to fully understand and exceed their clients' expectations, compromising the overall quality of their services. As a result, client satisfaction remains a critical variable in providing excellent service, which require airlines to precisely identify and address these needs (Arpita et al., 2023).

## DATASET DESCRIPTION

This dataset contains reviews of the top 10 rated airlines obtained from the Kaggle website <https://www.kaggle.com/datasets/sujalsuthar/airlines-reviews/download?datasetVersionNumber=1>. The reviews cover various aspects of the flight experience, including seat comfort, staff service, food and beverages, inflight entertainment, value for money, and overall rating. The dataset can be used for sentiment analysis, customer satisfaction analysis, visualisations, machine learning and other similar tasks.

It consists of 8,365 rows and 17 columns, these columns include:

Column names	Description
Title	Title of the review
Name	The reviewer's name
Review date	Date when the review was made
Verified	Shows if the review is verified or not
Reviews	The context of the review
Type of traveller	Shows the type of traveller the reviewer is. There are 4 types- solo leisure, family leisure, couple leisure or business.
Month Flown	The month the flight was taken
Route	The route of the flight
Class	The class of service, is it business class, economy class, premium class or first class
Airline	The airline used by the reviewer
Seat comfort	The reviewer gives a rating score between 1-5 for this particular service
Staff service	The reviewer gives a rating score between 1-5 for this particular service
Food and beverages	The reviewer gives a rating score between 1-5 for this particular service
Inflight entertainment	The reviewer gives a rating score between 1-5 for this particular service
Value for money	The reviewer gives a rating score between 1-5 to confirm if they had value for their money
Overall ratings	The reviewer gives an overall rating score between 1-10 for each airline based on their services
Recommended	This shows if the customer will recommend the airline or not

Table 1. Dataset Description

The airlines reviewed in this dataset are:

1. Singapore Airlines
2. Qatar Airways

3. All Nippon Airways
4. Emirates
5. Japan Airlines
6. Turkish Airlines
7. Air France
8. Cathay Pacific Airways
9. EVA Air
10. Korean Air

The data types represented in the dataset are integer, string and Boolean.

## **METHODOLOGY**

The dataset includes qualitative and quantitative data that will be used for the analysis to make useful predictions. The design is experimental, and it will be done using machine learning models, the research methods adopted in this project are stated below:

### **4.1 Data Collection**

The dataset is obtained from Kaggle on this link  
<https://www.kaggle.com/datasets/sujalsuthar/airlines-reviews/download?datasetVersionNumber=1>.

### **4.2 Data Preprocessing**

The processing stage includes the process of cleaning the raw data by resolving issues such as class imbalance, checking for missing values, standardising your dataset, feature selection (removing columns that may not be required for the analysis), and feature encoding. All of these will be implemented before the application of the machine learning algorithm.

## 4.3 Application of Machine Learning Algorithms

The dataset will be split into two, a test set and a training set. The machine model will be trained using the training dataset. The following methods will be used for this project:

### 4.3.1 Logistics Regression:

Kanade (2022) defined logistic regression as a supervised machine learning technique that performs binary classification tasks by predicting the likelihood of an outcome, occurrence, or observation. The model produces a binary outcome with only two potential values: yes/no, 0/1, or true/false.

Logical regression examines the relationship between one or more independent variables and assigns data to discrete classes. It is often used in predictive modelling, where the model calculates the mathematical probability of whether an instance falls into a particular group or another (Kanade, 2022).

### 4.3.2 Random Forest:

Random Forest is a widely used machine learning method that combines the outputs of different decision trees to get a conclusion. Because of its ease of use and flexibility, it has boosted its recognition, as it can handle classification and regression problems.

The random forest method is an extension of the bagging approach by joining bagging and feature randomness to create an uncorrelated forest of decision trees. Feature randomization, which can be called feature bagging or "the random subspace method," creates a random selection of features to ensure low correlation among decision trees. This is one important difference between decision trees and random forests. Random forests select only a subset of the available feature splits, whereas decision trees consider all of them (*What Is Random Forest?* / IBM, n.d.).

### 4.3.3 Decision Tree:

A decision tree is a non-parametric supervised learning method that is used for classification and regression tasks. It has a hierarchical tree structure that consists of a root node, branches, internal nodes, and leaf nodes.

A decision tree begins with a root node that has no incoming branches. The root node's outgoing



branches feed into internal nodes, known as decision nodes. Based on the available attributes, both node types are evaluated to generate homogeneous subsets and they are denoted by leaf nodes or terminal nodes. The leaf nodes reflect every possible outcome in the dataset (*What Is a Decision Tree?* / IBM, n.d.-b).

#### **4.4 Software Installation**

The required version of Hadoop and Spark was installed on Ubuntu. The proof of the installations can be seen in Fig 1, Fig 2 and Fig 3. Tableau software was also installed for visualization (Fig 4).

### **EXPERIMENTAL SECTION**

To begin this process PySpark and findspark and imported were installed on the Jupyter environment. The installation process can be seen in Fig 15. With the software being installed the next stage was to import all the necessary libraries including the machine learning libraries from PySpark (Fig 16). After this, the next step was to import that airline review dataset. The dataset was loaded into Spark DataFrame and then started the preprocessing stage (Fig 17).

#### **5.1 Preprocessing of the dataset**

In this stage, steps to prepare the data for the models are being carried out. This process involves cleaning the data to be able to fit into the language the machine understands and also removing some information like columns that may not be needed for the analysis.

The preprocessing stage begins with checking for missing or null values and dropping these columns from the DataFrame as seen in Fig 18. In the case of the airline review dataset, there were no missing values hence this step wasn't required. Fig 18 also shows the number of rows and columns in the dataset. The next step was to drop columns that were not required for the analysis, the experiment uses the ratings from these customers to predict if they will recommend that airline or not. Fig 19 shows the columns that were not needed for the experiment and the datatype of the rest of the columns. Notice that Pyspark is reading all of these columns as a string. To change it to its correct dtype we imported the datatype library from spark and used the cast() function to make these changes (Fig 20 and Fig 21). After this, we converted the target

variable column from string yes and no to integers 1 and 0 and the Boolean column-true and false to integers 1 and 0. After converting the datatypes and changing the target and Boolean column we noticed there were missing values and we dropped them, this reduced the number of rows (Fig 22, Fig 23).

The next step is to check for class imbalance, fig 22 shows that our target variable may not have been perfectly balanced however the degree of distribution between the two classes can be negligible.

Class imbalance is when some classes have more instances than others. However, in most cases, if the degree of distribution between the classes is minute it is usually neglected and no class balance technique is applied.

The next step was encoding categorical columns by converting them to binary to enable the machine to work with them as the machine understands binary or numeric columns. First, we pick out the categorical columns that require modification. Every unique string value of these columns is assigned to a unique numerical index to enable them to be useable for machine learning. Then encoding takes place to convert these numeric indices to binary vectors, next a pipeline is constructed to effectively apply the indexing and encoding processes. The pipeline transforms the columns of the original dataset to encode, this process can be seen in Fig 24. After the encoding process, we drop the original columns and use the encoded columns.

The next step is to separate the target variable and then combine all the numerical columns into one output column as in Fig 25, after this is the data scaling or normalization, this is done to have all the values in the data in one scale which is important for machine learning models as it can help enhance model accuracy (Fig 26).

The next step is to split our data into a testing set and a training set. 70% of the data is assigned to the training set and 30% of the data is assigned to the testing set. Setting seed is done to ensure that the same data is been used anytime the data is run again, it helps with consistency every time the data is run (Fig 27).

Once the preprocessing stage is done the data is now ready for training.

## 5.2 Model Training and Performance evaluation using Accuracy score.

In this stage, three models are implemented and used for this experiment. After training, the models are used to predict the outcomes of the testing data. The accuracy score will be used to check the model that predicts best.

Fig 28 shows the implementation of the logistic regression method.

Fig 29 shows the implementation of the random forest method.

Fig 30 shows the implementation of the decision method.

## 5.3 Tableau Visualization

The dataset was imported into the Tableau environment then I dropped a row with a value zero as the ratings for each service were between 1 and 5 and the overall rating was between 1 and 10, next I removed some columns that won't be used for the visualization process. This column includes title, name, review date, verified, and reviews. In the visualization process, we want to be able to check the airlines that had more positive reviews, the type of travel customers would prefer to make, the airline with the best services and the airline with the highest recommendation.

Fig 5 shows the airline with the highest overall rating is Qatar Airways, this shows that Qatar Airways had a lot of ratings and the airline with the least highest overall rating is Korean Air. The airline with the overall lowest rating is Turkish Airlines.

From Fig 6 we can conclude that the customers who went for solo leisure gave the highest overall ratings and the customers who went for Business gave the lowest ratings.

Fig 7 shows that customers in economy class gave the highest overall ratings however majority gave the lowest ratings. The majority were not okay with the services they received from these airlines.

We will also see ratings based on each service they received from these airlines:

Fig 8 shows that Qatar Airways had comfortable seats as they had the highest positive review while Korean Air had the lowest positive review however Turkish Airlines had the highest negative review.

In Fig 9 we can see that Qatar Airways has the best staff service while Turkish Airways doesn't have excellent staff service.

Fig 10 shows that Customers don't really like the food and beverages served by Turkish Airlines however there is an equal number that likes it. While a lot of customers like the food and beverages served by Qatar Airways. From the plot, we can also see that Japan had the worst food and beverage ratings.

Fig 11 shows that customers who used Qatar Airways were very entertained while customers who used Korean Air were very entertained however a lot of customers were not pleased with the services of Turkish Airlines.

Fig 12 shows that a lot of customers who used Qatar Airways agreed that they had value for their money.

From the charts, we can see that Qatar Airways has the highest chance of being positively recommended by customers because they have shown excellent services based on the ratings from the airline review dataset, this is proven in Fig 13. Fig 14 shows that over the years Qatar Airways has always provided good services to their customers.

## RESULT DISCUSSION

The models were trained using 14 features out of 17 features. These 14 features were of utmost importance in the prediction of airline customers' recommendations. We use the accuracy scores to determine the performance of the machine learning algorithms applied. Below is the accuracy scores and other metrics evaluation for each of the models:

Models	Accuracy Score	Precision	Recall	F1 Score	AUC-ROC
Logistic Regression	71.44%	71.42%	71.44%	71.43	80.1%
Random Forest	91.83%	92.1%	91.8%	91.9%	97.72%
Decision Tree	91.6%	91.9%	91.6%	91.7%	94.5%

Table 2. Accuracy score for each machine learning model applied.

## CONCLUSION AND FUTURE WORK

From the visualization of the dataset, we were able to know the airline that will likely have the highest positive recommendations based on the reviews by airline customers, we can also see that many airlines need to improve their services. I will recommend they are more in touch with their customers, take their reviews very seriously and possibly ask these customers what they would do differently, customers might have ideas that can help the business. Qatar Airways may have had the highest positive ratings however they have some customers that disagree with their services, this airline can go through their reviews to see where they may have lacked and work on those areas.

Based on the result of our experiment Random Forest classifier and Decision Tree classifier passed the threshold of 80% however Random Forest is the most suitable models for predicting the outcome of airline customer recommendation choice however other models can also be applied to this experiment.

For future works, I recommend applying other machine learning models like deep learning or neural networks and using advanced techniques like hyperparameter tuning for optimization which can enhance prediction accuracy.

## SOCIAL IMPACT OF THIS PROJECT

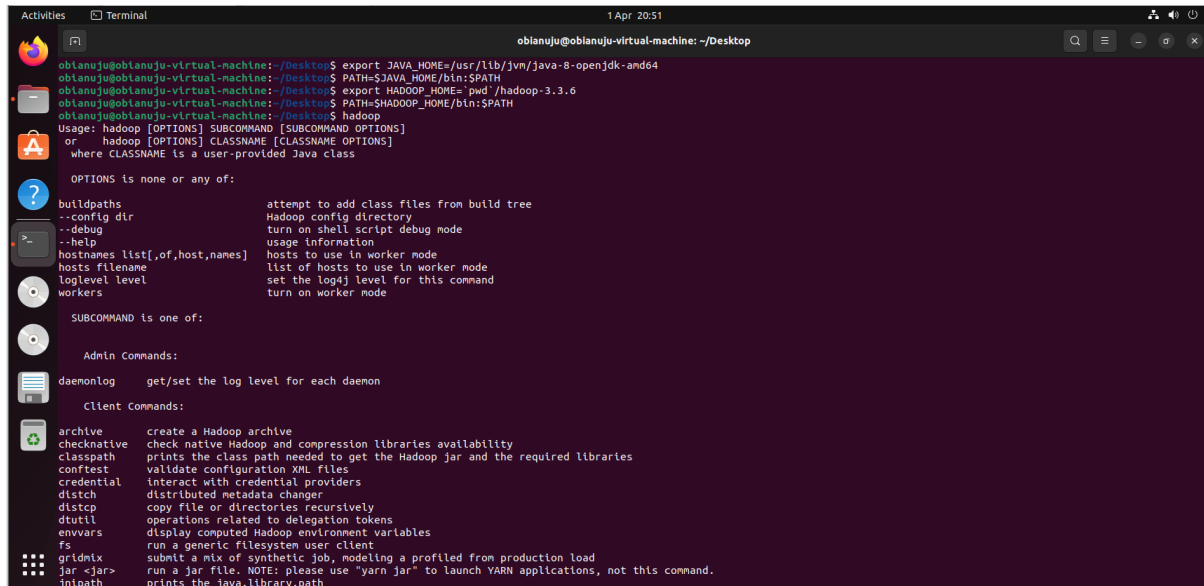
Being able to predict if a customer will come back to a business or recommend a business can help a business improve its services and using data insights and machine learning approaches can help businesses act early and faster as technology has made it easier to get insights faster.

These insights can provide an in-depth look at customers' attitudes and feelings towards various areas of the airline's services, allowing these airlines to better understand what drives their client satisfaction or disappointment and inform better decisions.

## REFERENCE

- Arpita, H. D., Ryan, A. A., Hossen, M. S., & Rahman, M. S. (2023). Airline Sentiments Unplugged: Leveraging Deep Learning for Customer Insights. *2023 5th International Conference on Sustainable Technologies for Industry 5.0 (STI)*. <https://doi.org/10.1109/sti59863.2023.10464925>
- Kanade, V. (2022, April 18). *Logistic Regression: equation, assumptions, types, and best practices* - Spiceworks. Spiceworks. <https://www.spiceworks.com/tech/artificial-intelligence/articles/what-is-logistic-regression/#:~:text=Logistic%20regression%20is%20a%20supervised%20machine%20learning%20algorithm%20that%20accomplishes,1%2C%20or%20true%2Ffalse>.
- Lee, S. M., & Lee, D. (2019). “Untact”: a new customer service strategy in the digital age. *Service Business*, 14(1), 1–22. <https://doi.org/10.1007/s11628-019-00408-2>
- Patel, A., Oza, P., & Agrawal, S. (2023). Sentiment Analysis of Customer Feedback and Reviews for Airline Services using Language Representation Model. *Procedia Computer Science*, 218, 2459–2467. <https://doi.org/10.1016/j.procs.2023.01.221>
- Savage, M. (2024, April 2). What is customer service in 2024: Definition, types, benefits, stats. *The Future of Commerce*. <https://www.the-future-of-commerce.com/2021/08/02/what-is-customer-service-definition-examples/>
- What is Random Forest? | IBM. (n.d.). <https://www.ibm.com/topics/random-forest>
- What is a Decision Tree? | IBM. (n.d.-b). <https://www.ibm.com/topics/decision-trees>

## APPENDIX I



```
oblanuju@oblanuju-virtual-machine: ~/Desktop$ export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-and64
oblanuju@oblanuju-virtual-machine: ~/Desktop$ PATH=$JAVA_HOME/bin:$PATH
oblanuju@oblanuju-virtual-machine: ~/Desktop$ export HADOOP_HOME=/usr/local/hadoop-3.3.6
oblanuju@oblanuju-virtual-machine: ~/Desktop$ PATH=$HADOOP_HOME/bin:$PATH
oblanuju@oblanuju-virtual-machine: ~/Desktop$ hadoop
Usage: hadoop [OPTIONS] SUBCOMMAND [SUBCOMMAND OPTIONS]
or hadoop [OPTIONS] CLASSNAME [CLASSNAME OPTIONS]
where CLASSNAME is a user-provided Java class

OPTIONS is none or any of:

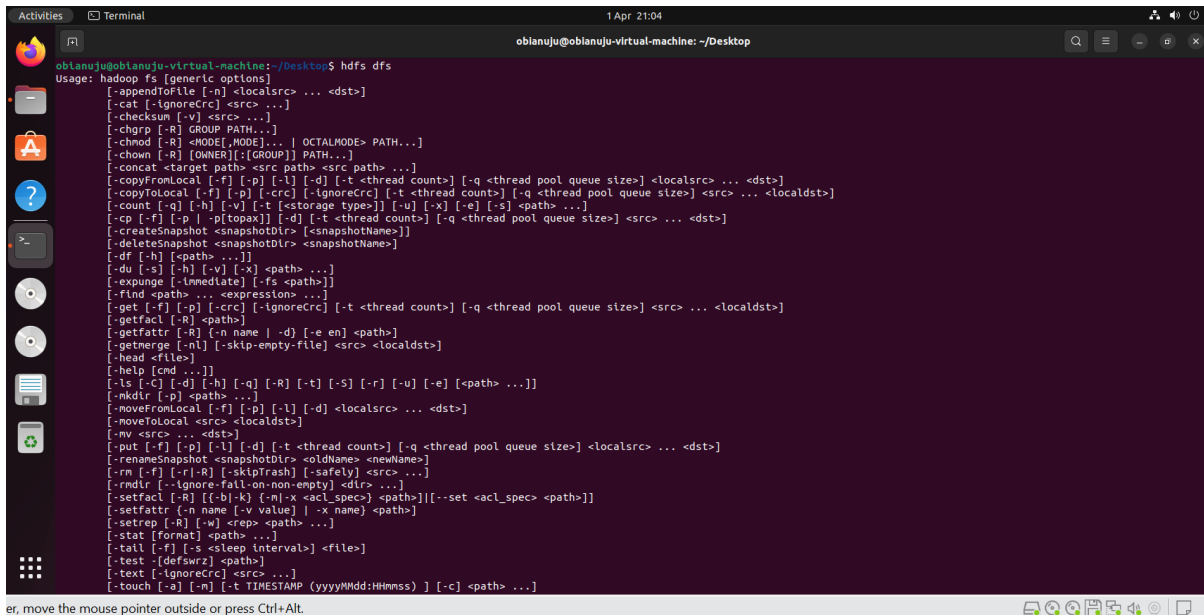
buildpaths      attempt to add class files from build tree
--config dir     Hadoop config directory
--debug          turn on shell script debug mode
--help           usage information
hostnames list[,of,host,names] hosts to use in worker mode
hosts filename   list of hosts to use in worker mode
loglevel level   set the log4j level for this command
workers          turn on worker mode

SUBCOMMAND is one of:

Admin Commands:
daemonlog        get/set the log level for each daemon

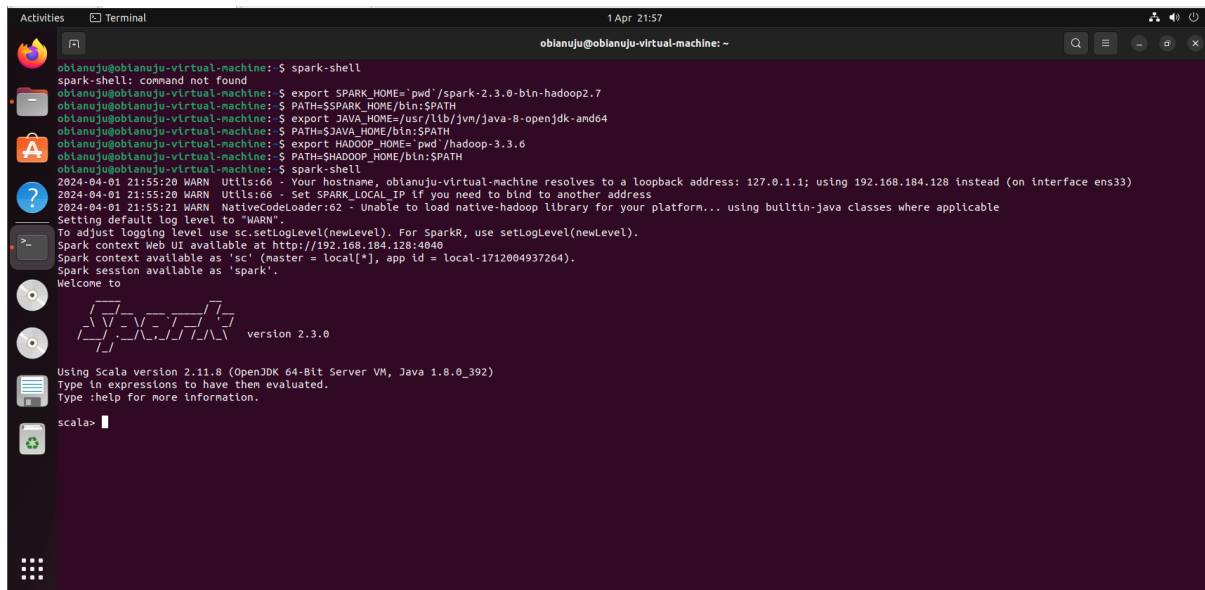
Client Commands:
archive          create a Hadoop archive
checknative      check native Hadoop and compression libraries availability
classpath        print the class path needed to get the Hadoop jar and the required libraries
confest         validate configuration XML files
credential       interact with credential providers
distch          distributed metadata changer
distcp          copy file or directories recursively
dutil           operations related to delegation tokens
envvars         display computed Hadoop environment variables
fs              run a generic filesystem user client
gridmix         submit a mix of synthetic job, modeling a profiled from production load
jar <jar>       run a jar file. NOTE: please use 'yarn jar' to launch YARN applications, not this command.
initpath        print the java.library.path
```

Fig 1. Proof of Hadoop being installed.



```
oblanuju@oblanuju-virtual-machine: ~/Desktop$ hdfs dfs
Usage: hadoop fs [generic options]
[-appendToFile [-n] <localsrc> ... <dst>]
[-cat [-ignoreCrc] <src> ...]
[-checksum [-v] <src> ...]
[-chgrp [-R] GROUP PATH...]
[-chmod [-R] <MODE>[[<MODE>]...] [-OCTALMODE] PATH...]
[-chown [-R] [OWNER][:[GROUP]] PATH...]
[-concat <target path> <src path> <src path> ...]
[-copyFromLocal [-f] [-p] [-l] [-d] [-t <thread count>] [-q <thread pool queue size>] <localsrc> ... <dst>]
[-copyToLocal [-f] [-p] [-rcrc] [-ignoreCrc] [-t <thread count>] [-q <thread pool queue size>] <src> ... <localdst>]
[-count [-q] [-h] [-v] [-t <storage type>]] [-u] [-c] [-e] [-s] <path> ...]
[-cp [-f] [-p] [-p[topax]] [-d] [-t <thread count>] [-q <thread pool queue size>] <src> ... <dst>]
[-createSnapshot <snapshotDir> [<snapshotName>]]
[-deleteSnapshot <snapshotDir> [<snapshotName>]]
[-df [-h] [<path> ...]]
[-du [-s] [-h] [-v] [-x] <path> ...]
[-expunge [-immediate] [-fs <path>]]
[-find <path> ... <expression> ...]
[-get [-f] [-p] [-rcrc] [-ignoreCrc] [-t <thread count>] [-q <thread pool queue size>] <src> ... <localdst>]
[-getfacl [-R] <path>]
[-getfattr [-R] [-n name | -d] [-e en] <path>]
[-getmerge [-nl] [-skip-empty-file] <src> <localdst>]
[-head <file>]
[-help [cmd ...]]
[-ls [-C] [-d] [-h] [-q] [-R] [-t] [-S] [-r] [-u] [-e] <path> ...]]
[-mkdir [-p] <path> ...]
[-moveFromLocal [-f] [-p] [-l] [-d] <localsrc> ... <dst>]
[-moveToLocal <src> <localdst>]
[-mv <src> ... <dst>]
[-put [-f] [-p] [-l] [-d] [-t <thread count>] [-q <thread pool queue size>] <localsrc> ... <dst>]
[-renameSnapshot <snapshotDir> <oldName> <newName>]
[-rm [-f] [-r] [-R] [-skipTrash] [-safely] <src> ...]
[-rmrdir [-ignore-fail-on-non-empty] <dir> ...]
[-setfacl [-R] [[-b|-k] [-n] <x> <acl_spec>] <path>][[--set <acl_spec>] <path>]]
[-setfattr [-n name [-v value] | -x name] <path>]
[-setrep [-R] [-w] <rep> <path> ...]
[-stat [format] <path> ...]
[-tail [-f] [-s <sleep interval>] <file>]
[-test [-defswr] <path>]
[-text [-ignoreCrc] <src> ...]
[-touch [-a] [-m] [-t TIMESTAMP (yyyyMMdd:HHmmss) ] [-c] <path> ...]
```

Fig 2. Calling Hadoop File System (HDFS).



The terminal window shows the installation of Spark on a virtual machine. The user runs `spark-shell`, which fails with "command not found". They then set environment variables for `SPARK_HOME`, `PATH`, `JAVA_HOME`, and `HADOOP_HOME`. After running `spark-shell` again, several warning messages appear regarding hostname resolution and native Hadoop library loading. The Spark context is established, and the user is welcomed to the Scala REPL.

```
obianuju@obianuju-virtual-machine:~$ spark-shell
spark-shell: command not found
obianuju@obianuju-virtual-machine:~$ export SPARK_HOME='pwd'/spark-2.3.0-bin-hadoop2.7
obianuju@obianuju-virtual-machine:~$ PATH=$SPARK_HOME/bin:$PATH
obianuju@obianuju-virtual-machine:~$ export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
obianuju@obianuju-virtual-machine:~$ PATH=$JAVA_HOME/bin:$PATH
obianuju@obianuju-virtual-machine:~$ export HADOOP_HOME='pwd'/hadoop-3.3.6
obianuju@obianuju-virtual-machine:~$ PATH=$HADOOP_HOME/bin:$PATH
obianuju@obianuju-virtual-machine:~$ spark-shell
2024-04-01 21:55:20 WARN Utils:66 - Your hostname, obianuju-virtual-machine resolves to a loopback address: 127.0.1.1; using 192.168.184.128 instead (on interface ens33)
2024-04-01 21:55:21 WARN NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Spark context Web UI available at http://192.168.184.128:4040
Spark context available as 'sc' (master = local[*], app id = local-1712004937264).
Spark session available as 'spark'.
Welcome to
scala>
```

Fig 3. Spark Installation.

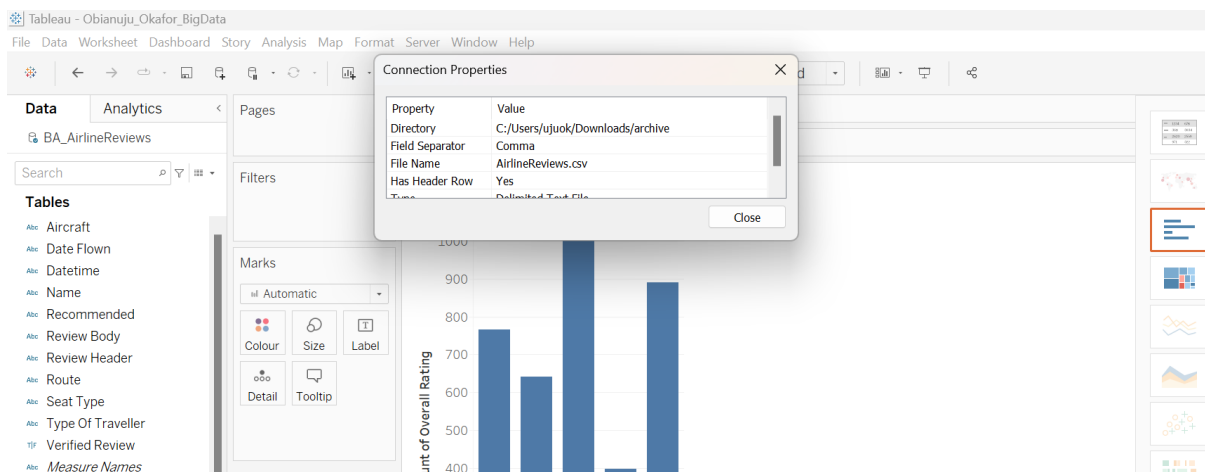


Fig 5. Tableau Proof



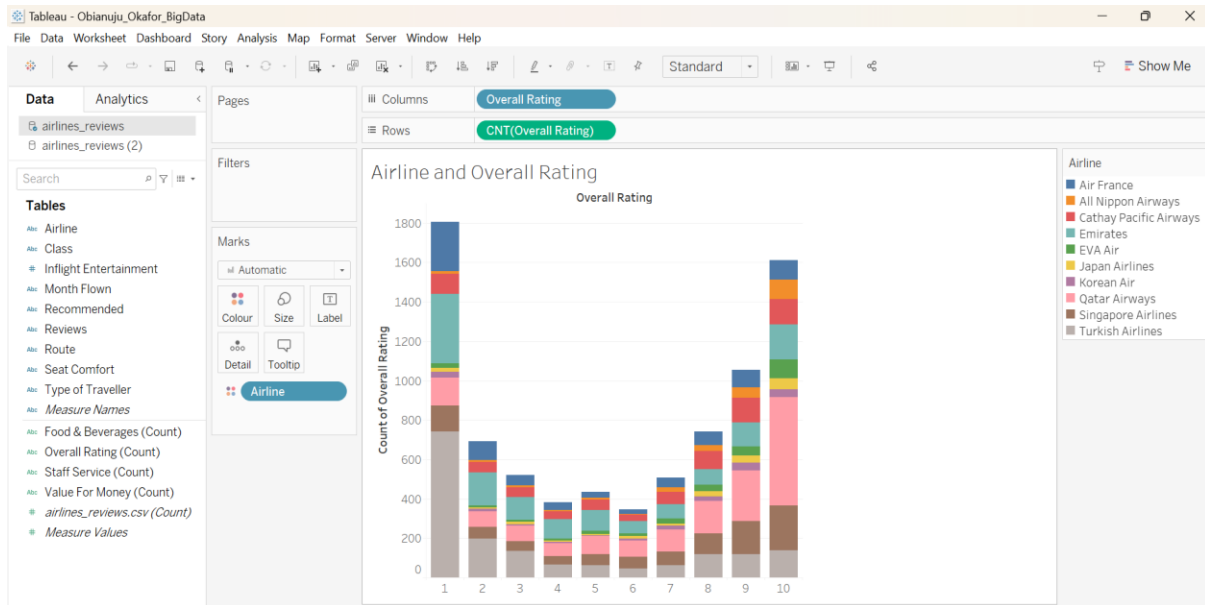
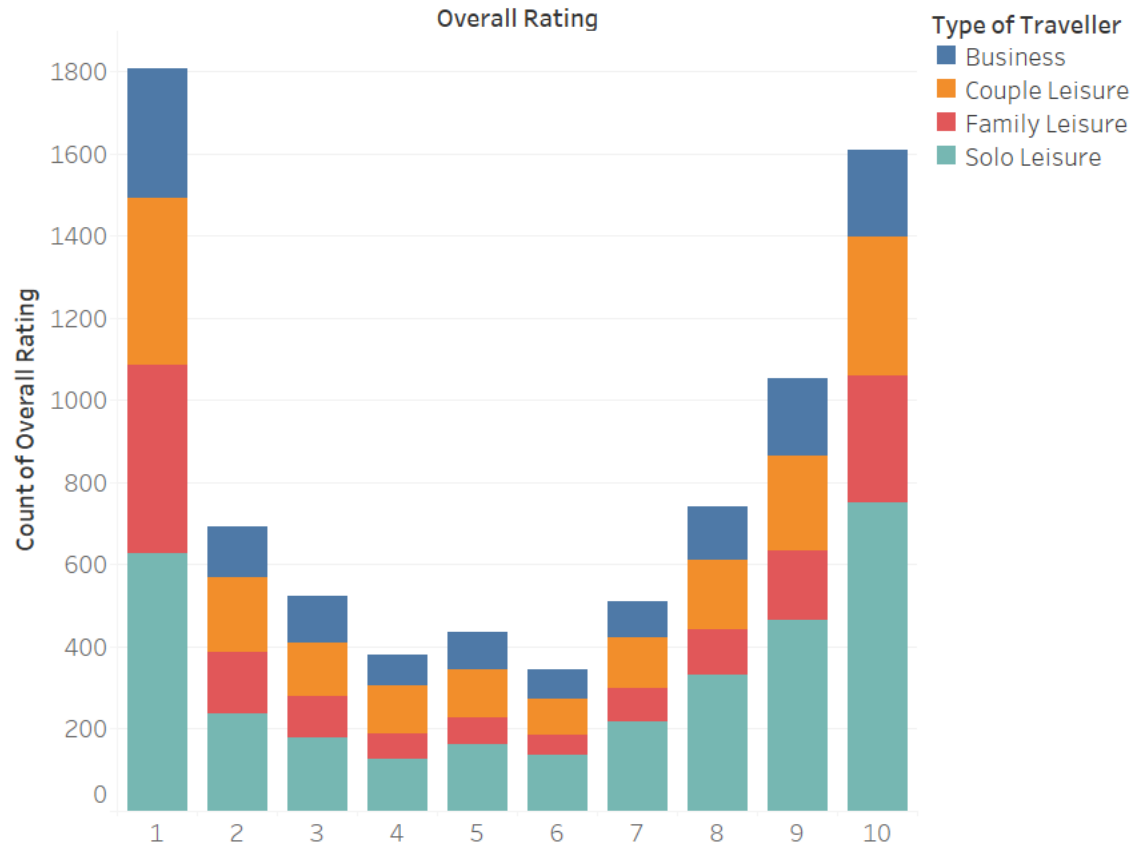


Fig 5. Overall ratings of each airline

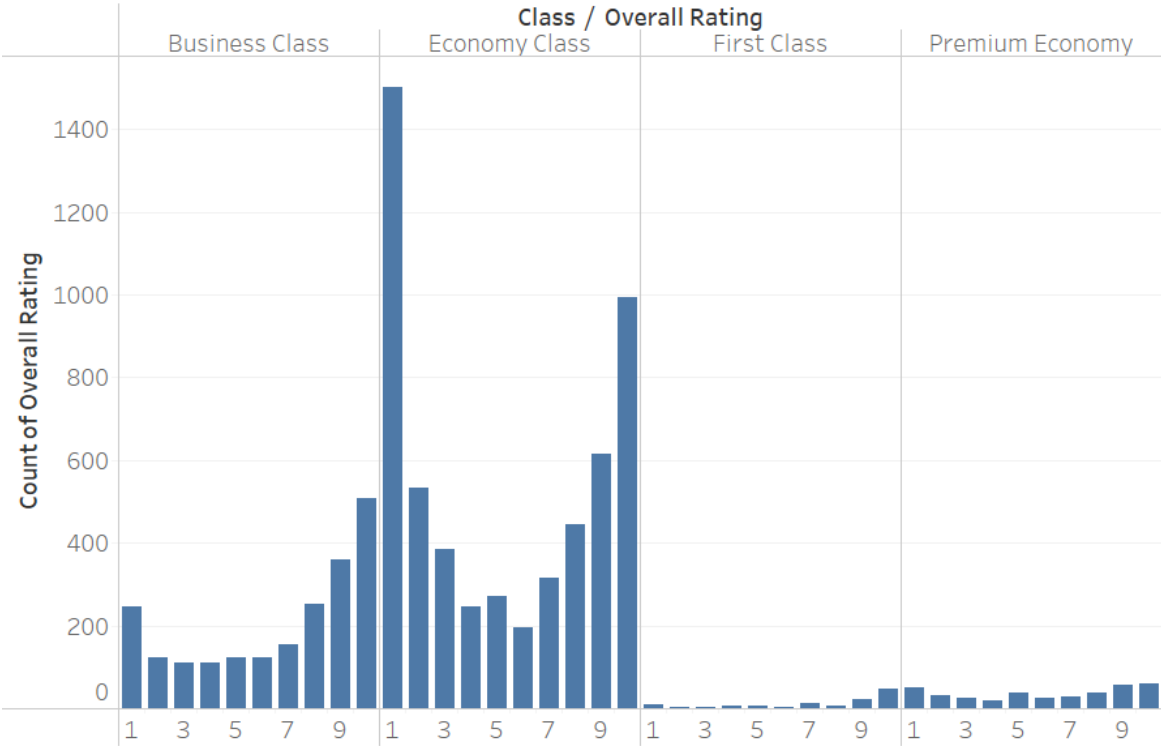
## Type of Traveller and Overall Rating



Count of Overall Rating for each Overall Rating. Colour shows details about Type of Traveller.

Fig 6. Overall ratings of type of traveller.

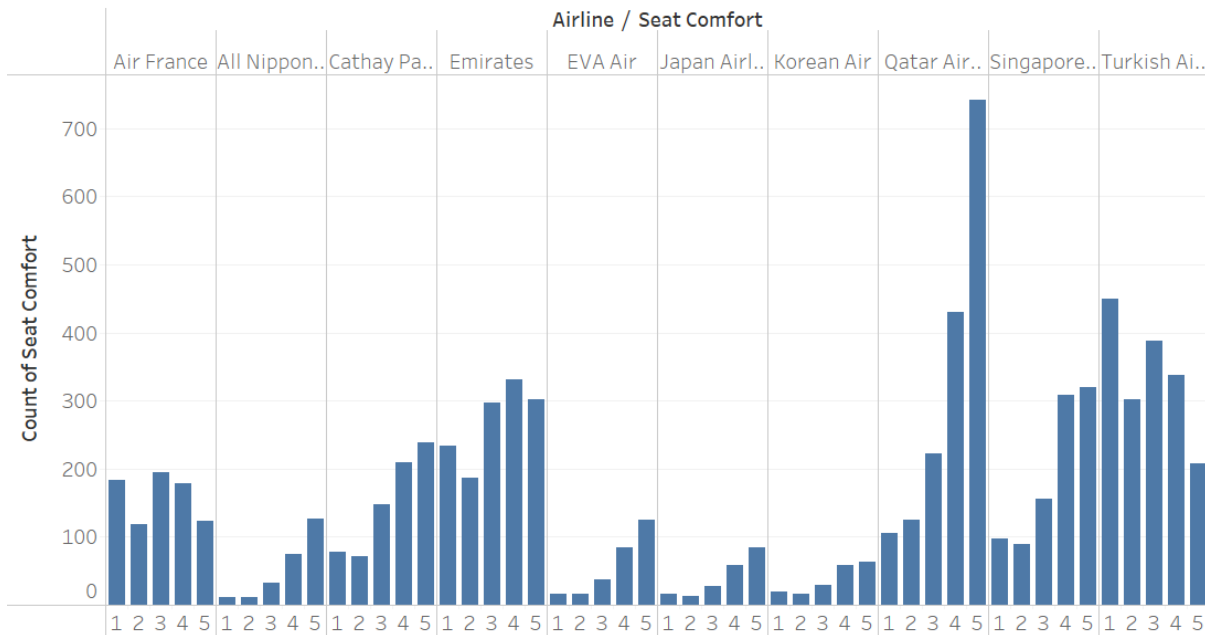
# Class and Overall Rating



Count of Overall Rating for each Overall Rating broken down by Class.

Fig 7 Overall ratings of Class

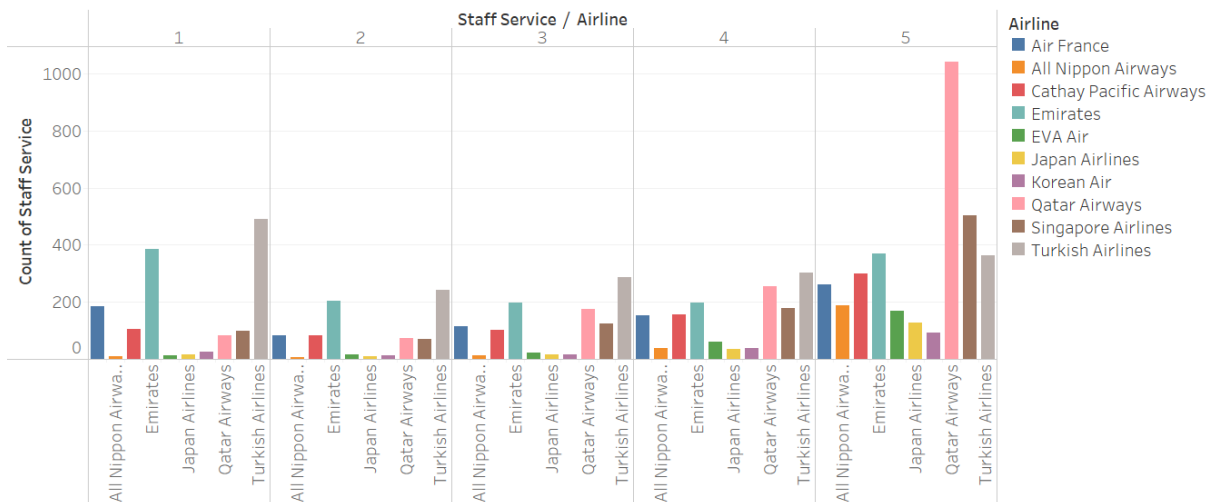
## Airline and Seat Comfort



Count of Seat Comfort for each Seat Comfort broken down by Airline.

Fig 8. Ratings of airlines and their seat comfort service

## Airline and Staff Service



Count of Staff Service for each Airline broken down by Staff Service. Colour shows details about Airline.

Fig 9. Ratings of airlines and their staff service

## Airline and Food and Beverage

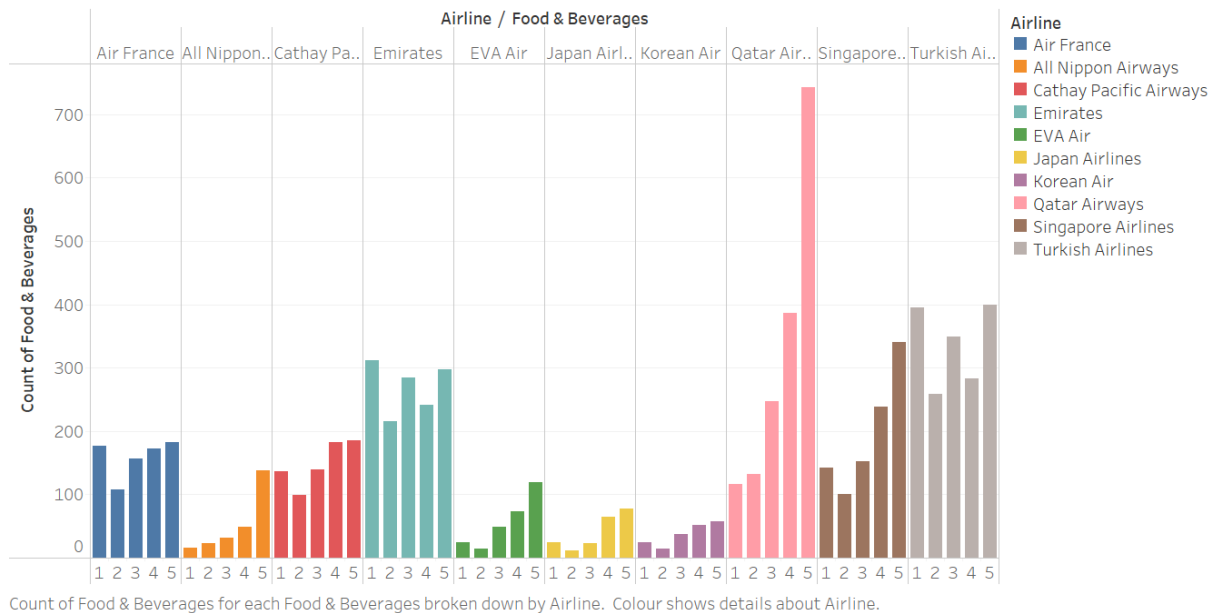
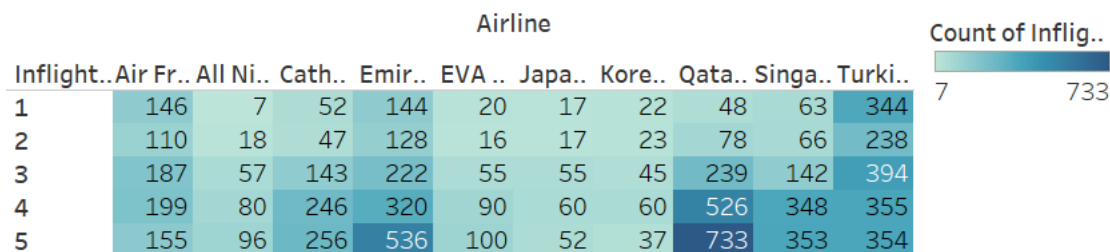


Fig 10. Ratings of airlines food and beverages

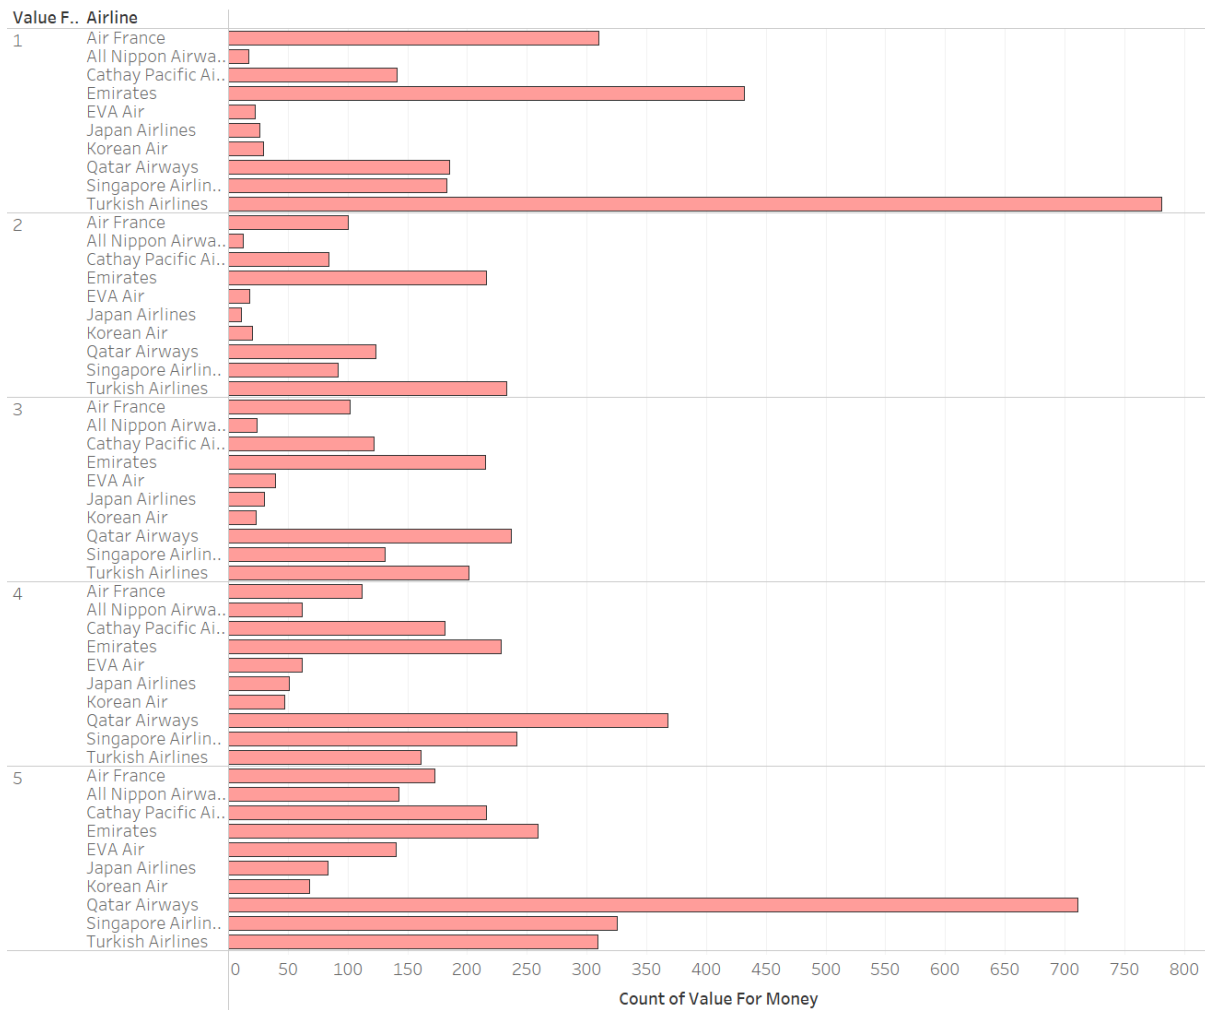
## Airline and Inflight Entertainment



Count of Inflight Entertainment broken down by Airline vs. Inflight Entertainment. Colour shows count of Inflight Entertainment. The marks are labelled by count of Inflight Entertainment.

Fig 11. Airlines' inflight entertainment ratings

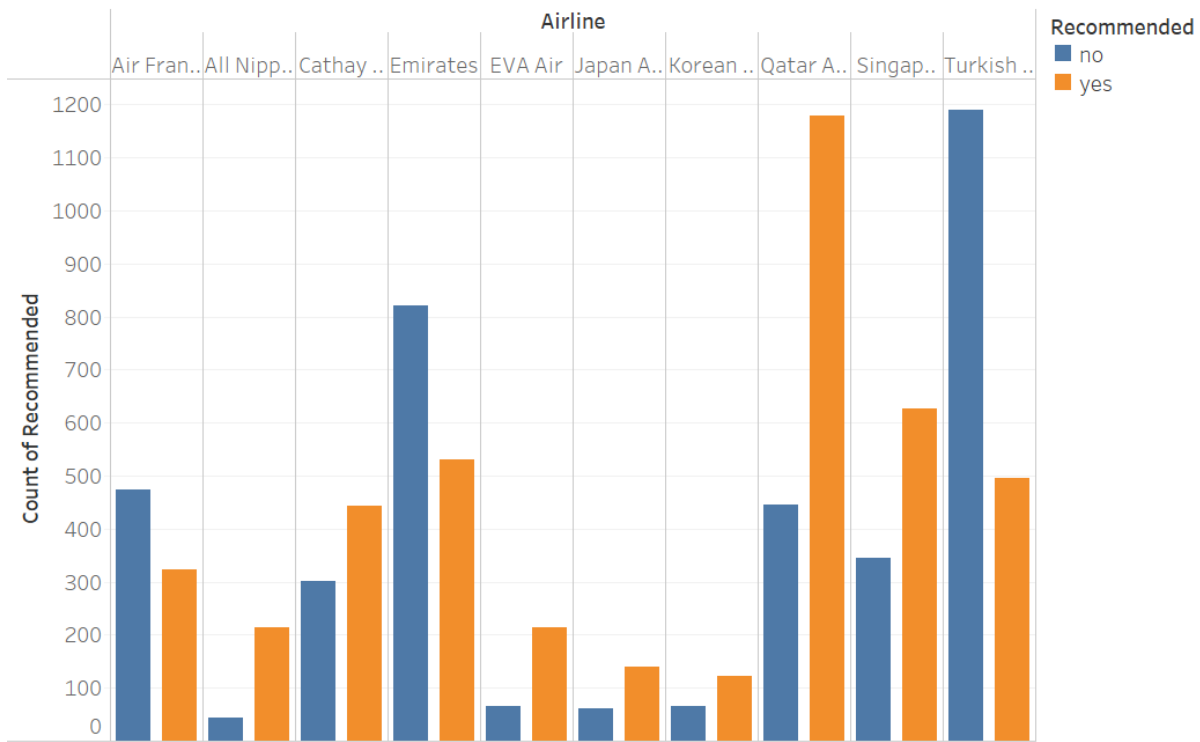
## Airline and Value for Money



Count of Value For Money for each Airline broken down by Value For Money.

Fig 12. Airlines that gave value for the money been paid by their customers.

## Airline and Recommendation



Count of Recommended for each Recommended broken down by Airline. Colour shows details about Recommended.

Fig 13. Airlines that customers will or will not recommend.

# Month Flown, Airline and Recommendation

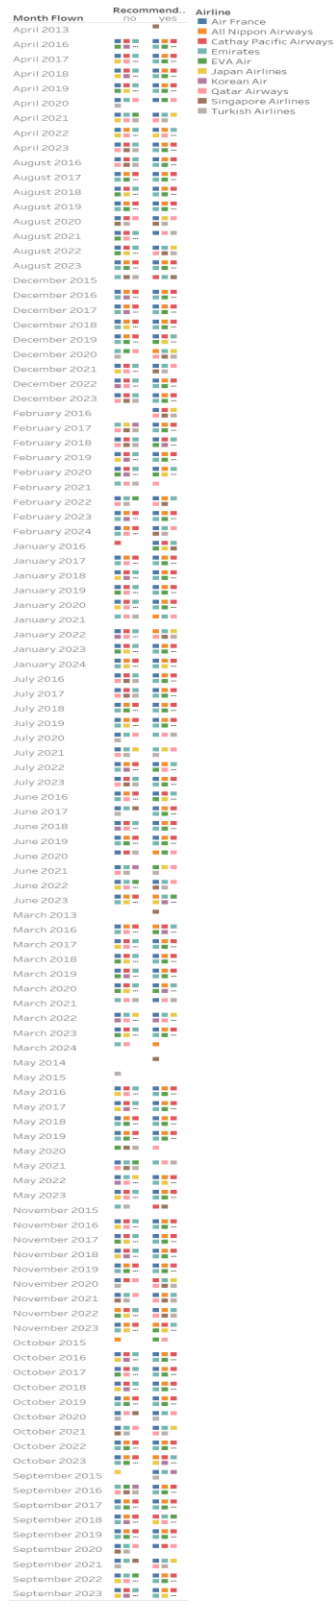
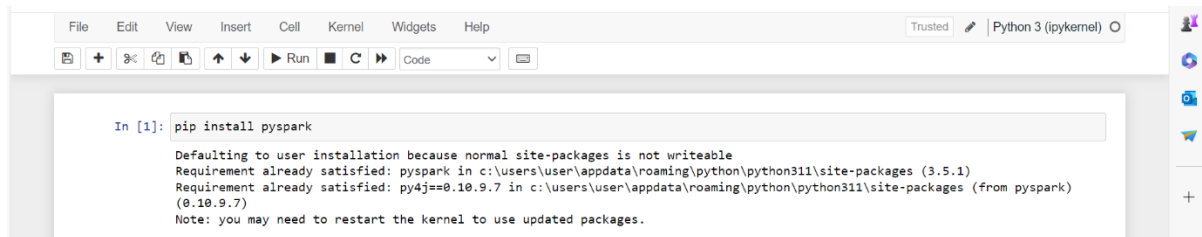




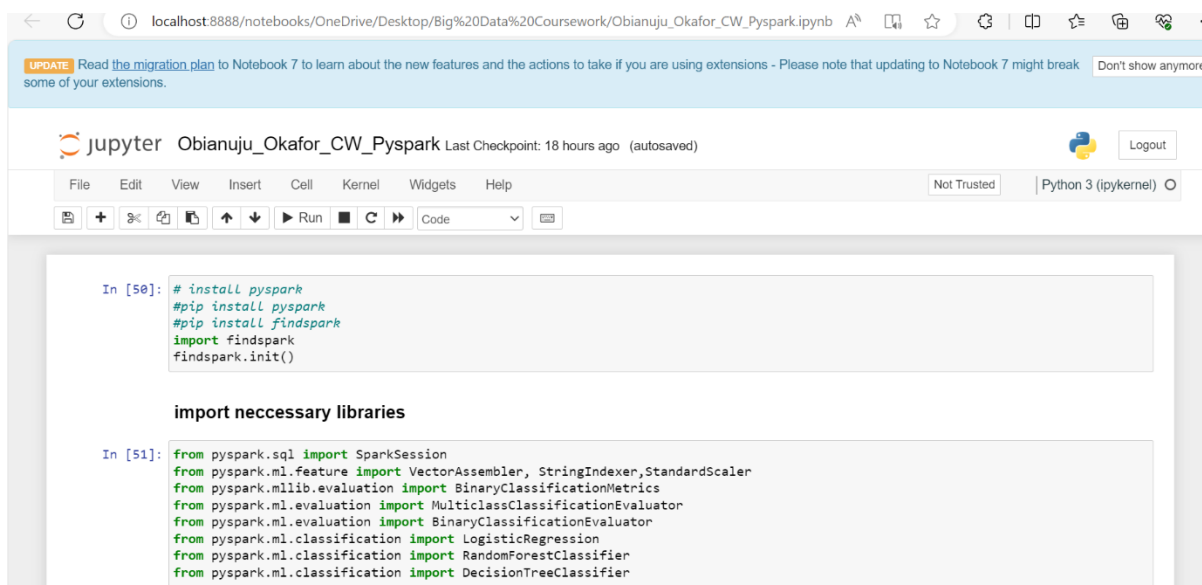
Fig 14. Airlines' recommendations from 2013 to 2024



```
In [1]: pip install pyspark

Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: pyspark in c:\users\User\appdata\roaming\python\python311\site-packages (3.5.1)
Requirement already satisfied: py4j==0.10.9.7 in c:\users\User\appdata\roaming\python\python311\site-packages (from pyspark) (0.10.9.7)
Note: you may need to restart the kernel to use updated packages.
```

Fig 15. Pyspark Installation.



```
In [50]: # install pyspark
#pip install pyspark
#pip install findspark
import findspark
findspark.init()

import necessary libraries

In [51]: from pyspark.sql import SparkSession
from pyspark.ml.feature import VectorAssembler, StringIndexer, StandardScaler
from pyspark.ml.evaluation import BinaryClassificationMetrics
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.classification import LogisticRegression
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.classification import DecisionTreeClassifier
```

Fig 16. Importing findspark and all the libraries needed.

```
In [98]: # Create a SparkSession
spark = SparkSession.builder.appName("AirlineRecommendationPredictor").master("local").getOrCreate()

In [99]: # Load the dataset
data = spark.read.csv("C:/Users/ujuok/OneDrive/Desktop/Big Data Coursework/airlines_reviews.csv", header=True, inferSchema=True)
```

Fig 17. Load dataset in Pyspark.

```
In [100]: import pyspark.sql.functions as fc
print((data.count(), len(data.columns)))
#data.describe().show()

# finding null values in each column
data_null = data.agg(*[fc.count(fc.when(fc.isNull(c), c)).alias(c) for c in data.columns])
#data_null.show()
# no null values

(8365, 17)
```

Fig 18. Number of rows and columns of the dataset and checking for missing values.

```
In [101]: # Drop columns not necessary for classification
airline_data = data.drop('Name', 'Review Date', 'Month Flown')

In [102]: airline_data.dtypes
#airline_data.summary
#airline_data.describe().show()

Out[102]: [('Title', 'string'),
('Airline', 'string'),
('Verified', 'string'),
('Reviews', 'string'),
('Type of Traveller', 'string'),
('Route', 'string'),
('Class', 'string'),
('Seat Comfort', 'string'),
('Staff Service', 'string'),
('Food & Beverages', 'string'),
('Inflight Entertainment', 'string'),
('Value For Money', 'string'),
('Overall Rating', 'string'),
('Recommended', 'string')]
```

Fig 19. Column drop and column dtype.

```
In [103]: # converting the interger columns to intergers

from pyspark.sql.types import IntegerType
from pyspark.sql.types import BooleanType

airline_data = airline_data.withColumn("Seat Comfort", airline_data["Seat Comfort"].cast(IntegerType()))
airline_data = airline_data.withColumn("Staff Service", airline_data["Staff Service"].cast(IntegerType()))
airline_data = airline_data.withColumn("Food & Beverages", airline_data["Food & Beverages"].cast(IntegerType()))
airline_data = airline_data.withColumn("Inflight Entertainment", airline_data["Inflight Entertainment"].cast(IntegerType()))
airline_data = airline_data.withColumn("Value For Money", airline_data["Value For Money"].cast(IntegerType()))
airline_data = airline_data.withColumn("Overall Rating", airline_data["Overall Rating"].cast(IntegerType()))
airline_data = airline_data.withColumn("Verified", airline_data["Verified"].cast(BooleanType()))
```

Fig 20. Converting the columns to the correct datatype.

```
In [104]: airline_data.dtypes
Out[104]: [('Title', 'string'),
('Airline', 'string'),
('Verified', 'boolean'),
('Reviews', 'string'),
('Type of Traveller', 'string'),
('Route', 'string'),
('Class', 'string'),
('Seat Comfort', 'int'),
('Staff Service', 'int'),
('Food & Beverages', 'int'),
('Inflight Entertainment', 'int'),
('Value For Money', 'int'),
('Overall Rating', 'int'),
('Recommended', 'string')]
```

Fig 21. Converted columns.

Jupyter Obianuju\_Okafor\_CW\_Pyspark Last Checkpoint: Yesterday at 02:02 (autosaved)

```
In [105]: # Converting the target variable column from string no and yes to integer 0 and 1
from pyspark.sql.functions import when

airline_data = airline_data.withColumn('Recommended_new', when(airline_data.Recommended=='yes', 1).otherwise(0))
airline_data = airline_data.drop("Recommended")

In [106]: # Converting the boolean variable column from bool False and True to integer 0 and 1
from pyspark.sql.functions import when

airline_data = airline_data.withColumn('Verified_new', when(airline_data.Verified==True, 1).otherwise(0))
airline_data = airline_data.drop("Verified")

In [107]: airline_data=airline_data.dropna()

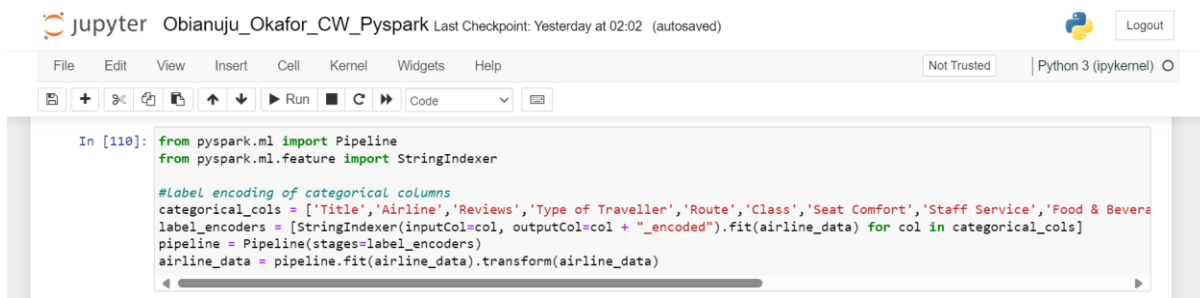
In [108]: airline_data.groupBy('Recommended_new').count().orderBy('count').show()

+-----+-----+
|Recommended_new|count|
+-----+-----+
|0|3408|
|1|3948|
+-----+-----+
```

Fig 22.

```
In [109]: print((airline_data.count(), len(airline_data.columns)))
(7356, 14)
```

Fig 23. The number of rows and columns after converting and dropping



```
In [110]: from pyspark.ml import Pipeline
from pyspark.ml.feature import StringIndexer

#Label encoding of categorical columns
categorical_cols = ['Title','Airline','Reviews','Type of Traveller','Route','Class','Seat Comfort','Staff Service','Food & Bever
label_encoders = [StringIndexer(inputCol=col, outputCol=col + "_encoded").fit(airline_data) for col in categorical_cols]
pipeline = Pipeline(stages=label_encoders)
airline_data = pipeline.fit(airline_data).transform(airline_data)
```

Fig 24. Label encoding of categorical variables.

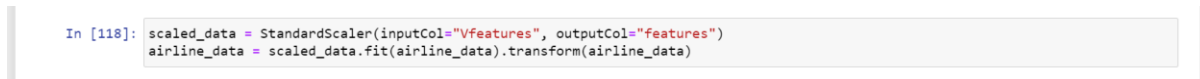


```
In [116]: # separating target variable from other features
features = airline_data.drop("Recommended_new")

# putting all the other features as one
features_col = features.columns
print(features_col)
assembler = VectorAssembler(inputCols=features_col, outputCol="Vfeatures")
airline_data = assembler.transform(airline_data)
#airline_data.show(2)
airline_data = airline_data.select("Vfeatures", "Recommended_new")

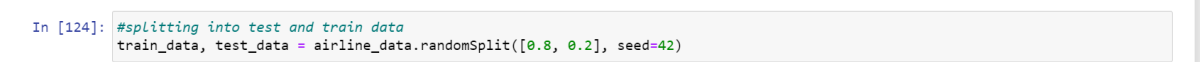
['Title_encoded', 'Airline_encoded', 'Reviews_encoded', 'Type of Traveller_encoded', 'Route_encoded', 'Class_encoded', 'Seat Co
mfort_encoded', 'Staff Service_encoded', 'Food & Beverages_encoded', 'Value For Money_encoded', 'Overall Rating_encoded', 'Infl
ight Entertainment_encoded', 'Verified_new_encoded']
```

Fig 25. Separating the target features from the others



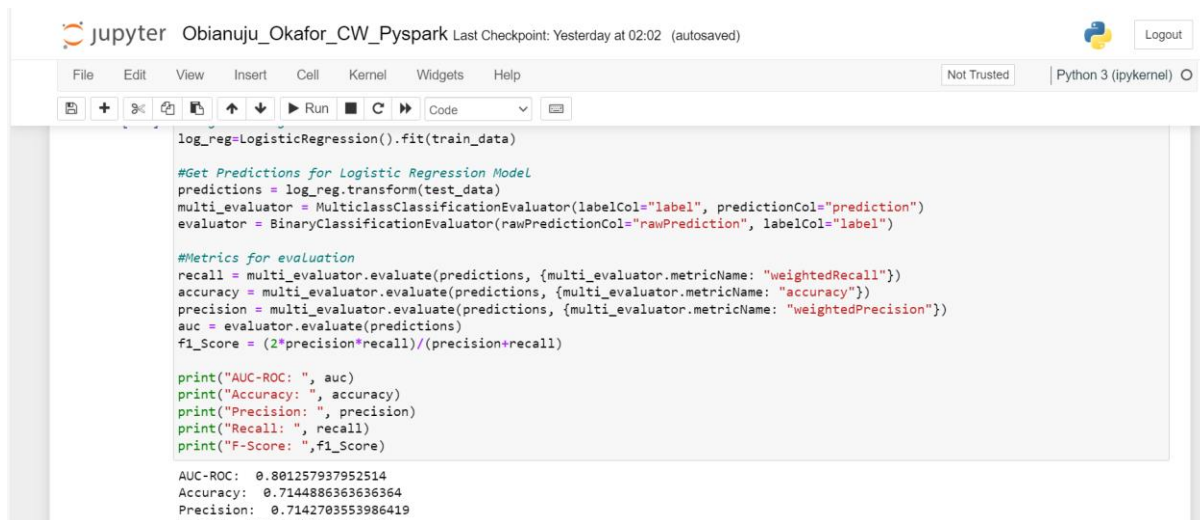
```
In [118]: scaled_data = StandardScaler(inputCol="Vfeatures", outputCol="features")
airline_data = scaled_data.fit(airline_data).transform(airline_data)
```

Fig 26. Scaling the data



```
In [124]: #splitting into test and train data
train_data, test_data = airline_data.randomSplit([0.8, 0.2], seed=42)
```

Fig 27. Splitting the data



The image shows a Jupyter Notebook interface with the title 'Obianuju\_Okafor\_CW\_Pyspark'. The last checkpoint is 'Yesterday at 02:02 (autosaved)'. The notebook contains a single code cell with the following Python code:

```
log_reg=LogisticRegression().fit(train_data)

#Get Predictions for Logistic Regression Model
predictions = log_reg.transform(test_data)
multi_evaluator = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction")
evaluator = BinaryClassificationEvaluator(rawPredictionCol="rawPrediction", labelCol="label")

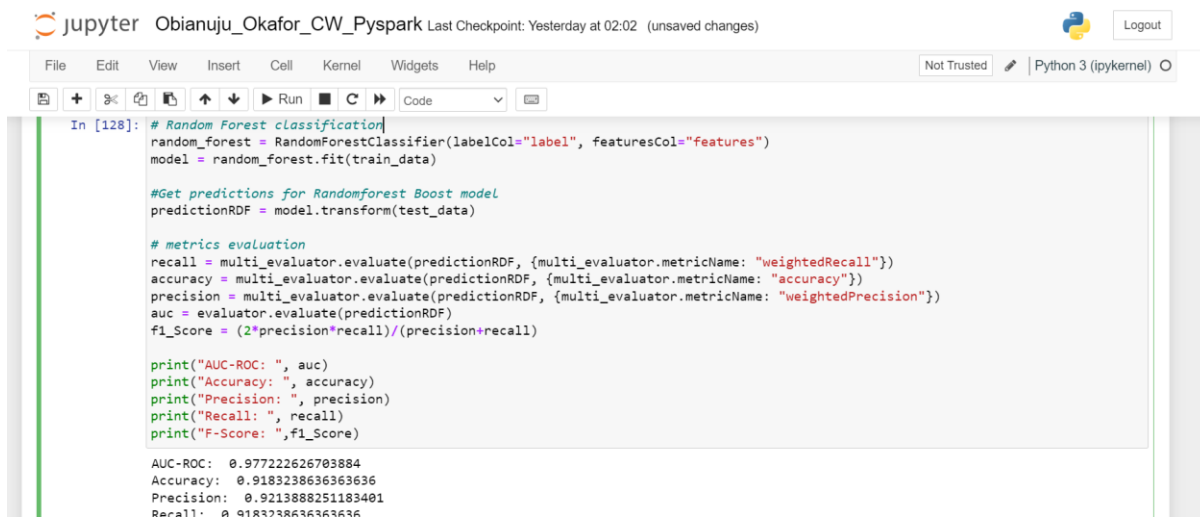
#Metrics for evaluation
recall = multi_evaluator.evaluate(predictions, {multi_evaluator.metricName: "weightedRecall"})
accuracy = multi_evaluator.evaluate(predictions, {multi_evaluator.metricName: "accuracy"})
precision = multi_evaluator.evaluate(predictions, {multi_evaluator.metricName: "weightedPrecision"})
auc = evaluator.evaluate(predictions)
f1_Score = (2*precision*recall)/(precision+recall)

print("AUC-ROC: ", auc)
print("Accuracy: ", accuracy)
print("Precision: ", precision)
print("Recall: ", recall)
print("F-Score: ", f1_Score)
```

The output of the code is displayed below the code cell:

```
AUC-ROC:  0.801257937952514
Accuracy: 0.7144886363636364
Precision: 0.7142703553986419
Recall:   0.7144886363636364
```

Fig 28. Logistic regression classification



The image shows a Jupyter Notebook interface with the title 'Obianuju\_Okafor\_CW\_Pyspark'. The last checkpoint is 'Yesterday at 02:02 (unsaved changes)'. The notebook contains a single code cell with the following Python code:

```
In [128]: # Random Forest classification
random_forest = RandomForestClassifier(labelCol="label", featuresCol="features")
model = random_forest.fit(train_data)

#Get predictions for Randomforest Boost model
predictionRDF = model.transform(test_data)

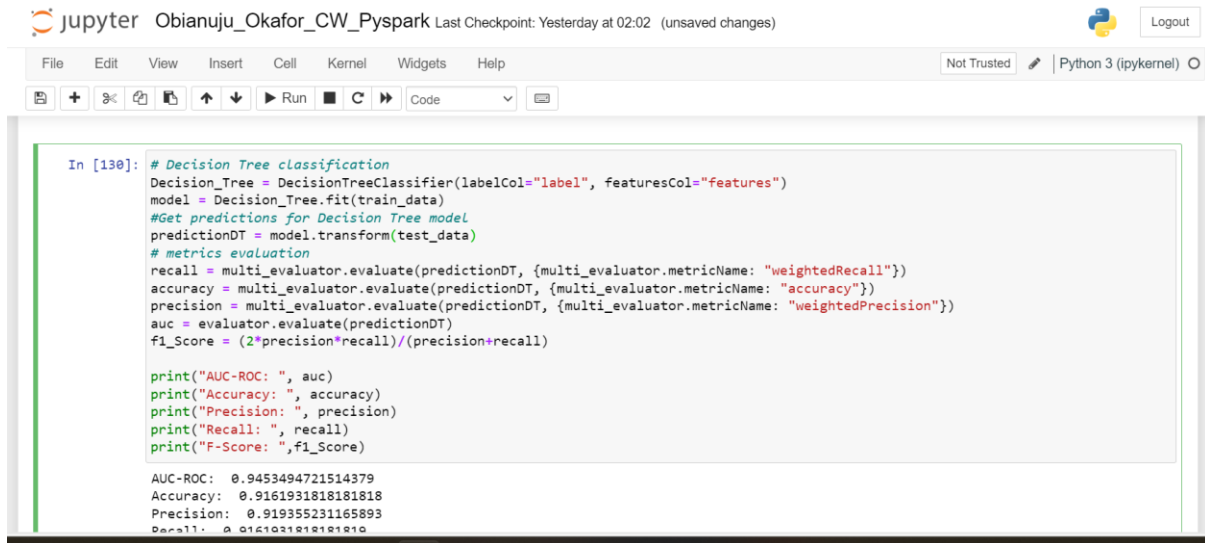
# metrics evaluation
recall = multi_evaluator.evaluate(predictionRDF, {multi_evaluator.metricName: "weightedRecall"})
accuracy = multi_evaluator.evaluate(predictionRDF, {multi_evaluator.metricName: "accuracy"})
precision = multi_evaluator.evaluate(predictionRDF, {multi_evaluator.metricName: "weightedPrecision"})
auc = evaluator.evaluate(predictionRDF)
f1_Score = (2*precision*recall)/(precision+recall)

print("AUC-ROC: ", auc)
print("Accuracy: ", accuracy)
print("Precision: ", precision)
print("Recall: ", recall)
print("F-Score: ", f1_Score)
```

The output of the code is displayed below the code cell:

```
AUC-ROC:  0.977222626703884
Accuracy: 0.9183238636363636
Precision: 0.9213888251183401
Recall:   0.9183238636363636
```

Fig 29. Random forest classification



The image shows a Jupyter Notebook interface with a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a toolbar. The notebook is titled 'Obianuju\_Okafor\_CW\_Pyspark' and shows the last checkpoint from yesterday at 02:02. The code in the cell is a Python script for a Decision Tree classification. It imports necessary libraries, fits a model on training data, makes predictions on test data, and evaluates the model using various metrics. The output of the code is printed to the console.

```
In [130]: # Decision Tree classification
Decision_Tree = DecisionTreeClassifier(labelCol="label", featuresCol="features")
model = Decision_Tree.fit(train_data)
#Get predictions for Decision Tree model
predictionDT = model.transform(test_data)
# metrics evaluation
recall = multi_evaluator.evaluate(predictionDT, {multi_evaluator.metricName: "weightedRecall"})
accuracy = multi_evaluator.evaluate(predictionDT, {multi_evaluator.metricName: "accuracy"})
precision = multi_evaluator.evaluate(predictionDT, {multi_evaluator.metricName: "weightedPrecision"})
auc = evaluator.evaluate(predictionDT)
f1_Score = (2*precision*recall)/(precision+recall)

print("AUC-ROC: ", auc)
print("Accuracy: ", accuracy)
print("Precision: ", precision)
print("Recall: ", recall)
print("F-Score: ", f1_Score)

AUC-ROC:  0.9453494721514379
Accuracy:  0.9161931818181818
Precision:  0.919355231165893
Recall:  0.9161931818181818
```

Fig 30. Decision Tree.

## APPENDIX II

Code used for the implementation of the experiment:

```
from pyspark.sql import SparkSession
from pyspark.ml.feature import VectorAssembler, StringIndexer, StandardScaler
from pyspark.mllib.evaluation import BinaryClassificationMetrics
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.classification import LogisticRegression
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.classification import DecisionTreeClassifier
```

```
# Create a SparkSession
spark =
SparkSession.builder.appName("AirlineRecommendationPredictor").master("local").getOrCreate()

# Load the dataset
data = spark.read.csv("C:/Users/ujuok/OneDrive/Desktop/Big Data Coursework/airlines_reviews.csv",
header=True, inferSchema=True)
```

```

import pyspark.sql.functions as fc
print((data.count(), len(data.columns)))
#data.describe().show()

# finding null values in each column
data_null = data.agg(*[fc.count(fc.when(fc.isNull(c), c)).alias(c) for c in data.columns])
#data_null.show()

# no null values

# Drop columns not necessary for classification
airline_data = data.drop('Name', 'Review Date', 'Month Flown')

airline_data.dtypes
#airline_data.summary
#airline_data.describe().show()

# converting the integer columns from string to integers

from pyspark.sql.types import IntegerType
from pyspark.sql.types import BooleanType

airline_data = airline_data.withColumn("Seat Comfort", airline_data["Seat
Comfort"].cast(IntegerType()))
airline_data = airline_data.withColumn("Staff Service", airline_data["Staff Service"].cast(IntegerType()))
airline_data = airline_data.withColumn("Food & Beverages", airline_data["Food &
Beverages"].cast(IntegerType()))
airline_data = airline_data.withColumn("Inflight Entertainment", airline_data["Inflight
Entertainment"].cast(IntegerType()))
airline_data = airline_data.withColumn("Value For Money", airline_data["Value For
Money"].cast(IntegerType()))
airline_data = airline_data.withColumn("Overall Rating", airline_data["Overall
Rating"].cast(IntegerType()))

```

```
airline_data = airline_data.withColumn("Verified", airline_data["Verified"].cast(BooleanType()))
```

```
airline_data.dtypes
```

```
# Converting the target variable column from string no and yes to integer 0 and 1
```

```
from pyspark.sql.functions import when
```

```
airline_data = airline_data.withColumn('Recommended_new', when(airline_data.Recommended=='yes',  
1).otherwise(0))
```

```
airline_data = airline_data.drop("Recommended")
```

```
# Converting the boolean variable column from bool False and True to integer 0 and 1
```

```
from pyspark.sql.functions import when
```

```
airline_data = airline_data.withColumn('Verified_new', when(airline_data.Verified=='True',  
1).otherwise(0))
```

```
airline_data = airline_data.drop("Verified")
```

```
airline_data.groupBy('Recommended_new').count().orderBy('count').show()
```

```
print((airline_data.count(), len(airline_data.columns)))
```

```
from pyspark.ml import Pipeline
```

```
from pyspark.ml.feature import StringIndexer
```

```
#label encoding of categorical columns
```

```
categorical_cols = ['Title','Airline','Reviews','Type of Traveller','Route','Class','Seat Comfort','Staff  
Service','Food & Beverages','Value For Money','Overall Rating','Inflight Entertainment','Verified_new']
```

```
label_encoders = [StringIndexer(inputCol=col, outputCol=col + "_encoded").fit(airline_data) for col in  
categorical_cols]
```

```
pipeline = Pipeline(stages=label_encoders)
```

```
airline_data = pipeline.fit(airline_data).transform(airline_data)
```

```
airline_data.dtypes
```

```
airline_data = airline_data.drop('Title','Airline','Reviews','Verified_new','Type of  
Traveller','Route','Class','Seat Comfort','Staff Service','Food & Beverages','Value For Money','Overall  
Rating','Inflight Entertainment','Recommended')
```



```

airline_data.dtypes

# seperating target variable from other features
features = airline_data.drop("Recommended_new")

# putting all the other features as one
features_col = features.columns
print(features_col)
assembler = VectorAssembler(inputCols=features_col, outputCol="Vfeatures")
airline_data = assembler.transform(airline_data)
#airline_data.show(2)
airline_data = airline_data.select("Vfeatures", "Recommended_new")

airline_data.show(5)

scaled_data = StandardScaler(inputCol="Vfeatures", outputCol="features")
airline_data = scaled_data.fit(airline_data).transform(airline_data)
#renaming the target column to label
airline_data = airline_data.select("features", "Recommended_new")
airline_data = airline_data.withColumnRenamed("Recommended_new", "label")
#splitting into test and train data
train_data, test_data = airline_data.randomSplit([0.8, 0.2], seed=42)

# Logistic Regression
log_reg=LogisticRegression().fit(train_data)

#Get Predictions for Logistic Regression Model
predictions = log_reg.transform(test_data)
multi_evaluator = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction")
evaluator = BinaryClassificationEvaluator(rawPredictionCol="rawPrediction", labelCol="label")

#Metrics for evaluation
recall = multi_evaluator.evaluate(predictions, {multi_evaluator.metricName: "weightedRecall"})
accuracy = multi_evaluator.evaluate(predictions, {multi_evaluator.metricName: "accuracy"})

```

```

precision = multi_evaluator.evaluate(predictions, {multi_evaluator.metricName: "weightedPrecision"})
auc = evaluator.evaluate(predictions)
f1_Score = (2*precision*recall)/(precision+recall)

print("AUC-ROC: ", auc)
print("Accuracy: ", accuracy)
print("Precision: ", precision)
print("Recall: ", recall)
print("F-Score: ",f1_Score)

#Display the Logistic Regresssion predictions
predictions.show()

# Random Forest classification
random_forest = RandomForestClassifier(labelCol="label", featuresCol="features")
model = random_forest.fit(train_data)

#Get predictions for Randomforest Boost model
predictionRDF = model.transform(test_data)

# metrics evaluation
recall = multi_evaluator.evaluate(predictionRDF, {multi_evaluator.metricName: "weightedRecall"})
accuracy = multi_evaluator.evaluate(predictionRDF, {multi_evaluator.metricName: "accuracy"})
precision = multi_evaluator.evaluate(predictionRDF, {multi_evaluator.metricName:
"weightedPrecision"})
auc = evaluator.evaluate(predictionRDF)
f1_Score = (2*precision*recall)/(precision+recall)

print("AUC-ROC: ", auc)
print("Accuracy: ", accuracy)
print("Precision: ", precision)
print("Recall: ", recall)
print("F-Score: ",f1_Score)

```

```

predictionRDF.show()

# Decision Tree classification
Decision_Tree = DecisionTreeClassifier(labelCol="label", featuresCol="features")
model = Decision_Tree.fit(train_data)
#Get predictions for Decision Tree model
predictionDT = model.transform(test_data)
# metrics evaluation
recall = multi_evaluator.evaluate(predictionDT, {multi_evaluator.metricName: "weightedRecall"})
accuracy = multi_evaluator.evaluate(predictionDT, {multi_evaluator.metricName: "accuracy"})
precision = multi_evaluator.evaluate(predictionDT, {multi_evaluator.metricName: "weightedPrecision"})
auc = evaluator.evaluate(predictionDT)
f1_Score = (2*precision*recall)/(precision+recall)

print("AUC-ROC: ", auc)
print("Accuracy: ", accuracy)
print("Precision: ", precision)
print("Recall: ", recall)
print("F-Score: ",f1_Score)

predictionDT.show()

```